Databasesystems 2
Forum: https://forum-db.informatik.uni-tuebingen.de/c/ss18-db2

Assignment 3 (08.05.2018)

Submission: Tuesday, 15.05.2018, 10:00 AM

1. [15 Points] **Decomposition and MonetDB**

   You are given the following table in `decompositon.sql`:

   t

id	a	b	c	d
1	6	'a'	1	'f'
2	7	'b'	2	'g'
3	8	'c'	1	'f'
4	9	'd'	2	'g'
5	0	'e'	1	'f'

   Load the table into your MonetDB database by loading `decomposition.sql` using

   `mclient -l sql <dbname> <path/to/decomposition.sql>`

   (a) Split table `t` into four tables where each table holds one of the columns of `t` and the `id`. Write down the result. The *decomposed* tables should look as follows:

   ta

id	a
...	...

   tb

id	b
...	...

   tc

id	c
...	...

   td

id	d
...	...

   (b) Find any **non-trivial functional dependencies** in table `t` and write them down.

   (c) Remove any unneeded rows in tables `ta`, ..., `td`. A row is unneeded if it can be removed and we are still able to restore table `t` using a SQL query. Provide the shortened tables `ta`, ..., `td`.

   (d) Formulate a MonetDB SQL query over the shortened tables `ta`, ..., `td` whose result is equal to `t`.

   (e) Explain why MonetDB does **not** apply this particular space-saving scheme to decompose a wide table. Briefly compare the scheme to MonetDB's method of reassembling wide tables from BATs.

2. [15 Points] **Page Layout in *PostgreSQL***

Load `documents.sql` to create a table `documents(title CHAR(4), doc TEXT)`. Rows in `documents` are – due to type `TEXT` of column `doc` – of variable length and hence can grow and exceed the free page space on `UPDATE`. How does *PostgreSQL* handle that?

Use the *PostgreSQL* extension `pageinspect`[1] to observe *PostgreSQL*'s behavior. **Hand in all queries you used** and describe your findings **briefly**. Proceed as follows:

(a) Make use of function

```
heap_page_item_attrs(get_raw_page('documents', ⟨page⟩), 'documents')
```

to inspect the organization of rows on all pages of table `documents`. Next to each row's header information (`lp`: *slotno*, `lp_off`: row pointer, `lp_len`: *row size*, `t_ctid`: row version-chain pointer[2]), the function extracts the raw data of all attributes in an array `t_attrs` of type `bytea[]`. Use function `convert(str BYTEA)→TEXT` provided in `documents.sql` to convert this attribute data to `TEXT`.

(b) Find out the *RID* (*page*, *slotno*) and the *row size* of the row containing `'doc1'` as well as the free space left on its page.

(c) Perform an `UPDATE` on `'doc1'` doubling the size of its `doc` column, thus exceeding the free page space. Does the *RID* still point to the same row? How are the pages and the physical location of the row data reorganized?

(d) Perform an `UPDATE` on `'doc2'` growing its *row size* to more than 8 kB. The new row cannot fit into any page – even an empty one. How does *PostgreSQL* cope with that?

   i. How does the *row size* of the new row compare to the size of the inserted `doc`-value?

   ii. Read chapter 66.2.1 "*Out-of-line, on-disk TOAST storage*"[3] of the *PostgreSQL* documentation. Explain in your own words, how *"sliced bread"*[4] relates to *PostgreSQL* in terms of our current problem.

   iii. Search the system catalog table `pg_class`[5] to find the `relname` of the TOAST table associated with table `documents`.

   iv. Query table `pg_toast.⟨TOAST_relname⟩` to find out, how many *chunks* (rows of the toast table) have been created to store your new `doc`-value.

---

[1] https://www.postgresql.org/docs/current/static/pageinspect.html
[2] See Slide 08 of Chapter 5 "Row Updates"
[3] https://www.postgresql.org/docs/current/static/storage-toast.html#STORAGE-TOAST-ONDISK
[4] https://en.wikipedia.org/wiki/Toast
[5] https://www.postgresql.org/docs/current/static/catalog-pg-class.html